



How to self-assess and audit application deployment in public cloud?

Pinja Koskinen
Vesa Simola

Masters thesis
XXXXX 2018
Cyber security
Master's degree programme in cyber security

Author(s) Pinja Koskinen, Simola, Vesa	Type of publication Masters thesis	Date Month Year
		Language of publication: English
	Number of pages	Permission for web publication: x
Title of publication Title possible subtitle		
Degree programme		
Supervisor(s) Last name, First name		
Assigned by		
Abstract		

Contents

1	Introduction	1
2	General overview of cloud services	2
2.1	Cloud terminology and concepts	3
2.1.1	Infrastructure-as-a-service, IaaS	4
2.1.2	Platform-as-a-service, PaaS	7
2.1.3	Software-as-a-service, SaaS	7
2.1.4	Differences between cloud types	8
2.1.5	Mixtures of the two, aka hybrid cloud	10
3	Security in cloud	11
3.1	Security in private cloud	11
3.2	Security in community cloud	12
3.3	Security in public cloud	12
3.4	Introduction to security aspects in cloud	12
3.5	Security aspects of application development in public cloud	12
3.6	Introduction to cloud security and application deployment	12
3.6.1	Source code control in cloud application development	12
3.6.2	Secure processes for updating application and library code	12
3.6.3	Secure application administration procedures	12
3.7	Security and availability, and what to prepare for	13
3.7.1	File sharing service had authentication bug	13
3.7.2	Outsourced datacenter	14
3.7.3	Outsourced server etc. infrastructure, either partially or wholly	14
3.7.4	Large cloud providers are certified, are they?	15
3.7.5	Cloud native applications: Maximize the benefits of cloud to protect data and services	15
3.7.6	Ready-to-deploy security improvements in cloud	16

3.8	Security and availability cons of cloud deployments	17
3.8.1	Service level agreements	17
3.8.2	Privacy implications of running application in cloud	18
3.8.3	Dependant to external 3rd party provider	18
3.8.4	Dependency to connectivity	18
3.8.5	Application security cannot be outsourced even in SaaS	22
3.8.6	System security cannot be outsourced in IaaS	22
4	Self-assessment of cloud security posture	22
4.1	What is self-assessment of security and why bother?	23
4.2	Existing assessment methods and proposed controls: cloud security alliance . . .	23
4.2.1	Application and interface security	23
4.2.2	Audit assurance and compliance	24
4.2.3	Business continuity management & operational resilience	25
4.2.4	Change control & configuration management	26
4.3	Risk management in cloud landscape	26
4.4	Self-assessment questionnaire	29
4.4.1	Main points of interest and reasoning behind assessment questions	29
5	Auditing application deployment in cloud	29
5.1	Classified data: to cloud or not?	29
5.2	Overview of audit criteria for cloud environment	29
6	Conclusions	29
6.1	Complexity and applicability of audit criteria	29
6.2	Review of assessment and criteria	29
6.3	Further developments	29
7	References	29

1 Introduction

This thesis is about coming up with the self-assessment and auditing mechanism for Finnish applications that are either already deployed in public cloud or for applications that are candidates for such deployments. By being seen as flexible, reliable and cost effective mean of application delivery it is no wonder that the usage of cloud computing has spread far and wide. Yet, there are usecases where public cloud platforms might not be optimal, feasible or, in some cases, allowed options for deployment. This paper was written to better understand the implications of cloud computing versus the requirements of certain official security, confidentiality, integrity and availability requirements, some of which are actual hard requirements set by Finnish officials and some that are more like recommendations and best practices. To make the matter even more complex there are several providers who sell public cloud. For the most part, the goods being offered are fairly similar in fashion or at least they provide the same basic building block upon to build ones application. Still, there are differences that could effect on the audit result that determines whether the particular application is a good candidate for cloud deployment, one example of such difference could be the various connectivity options for attaching the users to the cloud application, another example could be the layers provided by the services provider for building the defence in depth setting that is seen as a good practice. These small, but important variations may have dire consequences concerning the confidentiality, integrity or availability of the application and hence they need to be evaluated when determining if the application is a good candidate for cloud deployment at all and what provider or providers to select for the actual application deployment.

This thesis starts by building up the basis of what is cloud computing and what it is made of, what kind of different cloud service categories are available currently and we'll also discuss some of the deployment models of cloud computing. Some emphasis is also put on the application design required to better exploit the possibilities of cloud by means of using stateless micro services with containers and orchestration tools. Security aspects of each of the factors and what they bring into the table are considered from the classical CIA-perspective, eg. Confidentiality, Integrity and Availability. Last portions of this thesis consist of the self-assessment questionnaire and the audit criteria. These two can be used to assess application for cloud deployment and on the other hand to audit

2 General overview of cloud services

Cloud service is generally understood as a product consisting of outsourced IT infrastructure, software components and some means for the customer to manage the service. This results in scalable distributed compute resources that are accessed using the network where several factors like CPU, memory and storage can be adjusted according to current needs (Hausman, Kirk, Susan L. Cook, and Telmo Sampaio. "Chapter 1 - What is Cloud Computing?". Cloud Essentials: CompTIA Authorized Courseware for Exam CLO-001. Sybex. 2013. Books24x7. <<http://common.books24x7.com.ezproxy.com> (accessed September 25, 2018)). Also, cloud services tend to have scalability built-in in terms of having some form of pay-as-you-grow model, meaning that the customer can start off with some amount of capacity but increase the deployment size as their requirements change, same is true for down

scaling the capacity to avoid unused capacity overhead. Avoiding costs is seen as one of the major benefits of cloud computing as many of the costly bits and bobs of IT infrastructure and related personnel can in many cases be outsourced. This means things including, but not limited to:

- Datacenter facilities
- Electricity, cooling etc.
- Computer hardware
- Operating systems to run on the hardware
- Storage systems
- Internet facing connectivity

This makes cloud computing seem like an ideal option for certain use cases where company has no interest of hosting the infrastructure on their own. Instead of maintaining the above mentioned infrastructure, company can concentrate on they key business be it application development or the sales generated from their Internet shop. This is not always the case though, as we've discovered in later parts of this thesis.

2.1 Cloud terminology and concepts

Next we'll cover some of the generic concepts and basic terminology involved in cloud computing. At the time of writing the cloud services can be split into three rough main categories (What is cloud computing Microsoft Azure 2018), and we'll discuss each of the aforementioned categories separately bellow. It should be noted that the categories below are subject to change as the cloud computing evolves and also there will likely be overlap in real life scenarios. Example of such overlap could be PaaS service providing Kubernetes workflow automation while at the same time Kubernetes itself could be seen as a SaaS service. This overlap is natural result of the categories building on top of each other. We'll also touch the few options for hosting the the cloud, either as a completely outsourced service, inhouse and the mixture of the two. Security and availability aspects of the categories and their hosting options are discussed in detail later in this document.

2.1.1 Infrastructure-as-a-service, IaaS

Infrastructure as a service is the most basic service in the cloud landscape, it generally means an offering consisting of infrastructure, physical or virtual machines and other related resources like storage of images, networking and security features such as firewalls and load balancers and bundles of software. (Saara Suikkanen, 2013 s13) Important aspect of note is that while customer has limited or no control of underlying architecture that is used to provision the cloud based services, customer is still responsible for proper use and care of the cloud resources, for example the configuration of an application. (STAMOU, Aikaterini, 2014,s5)

This is the basic service where one rents computing capacity in a form of facilities that house hardware, virtual machines, storage capacity and network bandwidth. Operating systems and some core infrastructure services such as NFS storage or firewalls could be provided as part of the service. This category of cloud computing provides the basic pay-as-you-grow model in terms of giving the customer the opportunity to scale their number of virtual machines, storage capacity or pretty much any of the basic building blocks. The benefit of the IaaS cloud for the customer is that certain data center related activities can be abstracted and used from for example a web interface or an API. There is no need to manage all levels of the infrastructure anymore and administrative tasks can mostly focus on server side level like operating systems management and maintenance, and third party software maintenance. (Kavis, Michael. "Infrastructure as a Service". Architecting the Cloud: Design Decisions for Cloud Computing Service Models (SaaS, PaaS, and IaaS). John Wiley & Sons. 2014. Books24x7. <<http://common.books24x7.com.ezproxy.jamk.fi:2048/toc.aspx?bookid=62597>> (accessed September 25, 2018), Infrastructure as a Service) As in this type of a cloud service only the infrastructure is provided, all software related development and administration responsibilities are left to the customer (Kavis, Michael. "Infrastructure as a Service". Architecting the Cloud: Design Decisions for Cloud Computing Service Models (SaaS, PaaS, and IaaS). John Wiley & Sons. © 2014. Books24x7. <<http://common.books24x7.com.ezproxy.jamk.fi:2048/toc.aspx?bookid=62597>> (accessed September 25, 2018), Infrastructure as a Service). While this brings a lot of freedom on how to create the software and configuration, nothing is provided as a service by the cloud provider. It should also be noted that while service provider can have multiple geographical sites from which

they provide their service allowing distribution of the resources, none of this is benefitting the software stack without adding additional features to the software. To benefit from this the cloud awareness and scalability has to be done either manually or as a part of the software. Some of these limitations and distribution of responsibilities are discussed next.

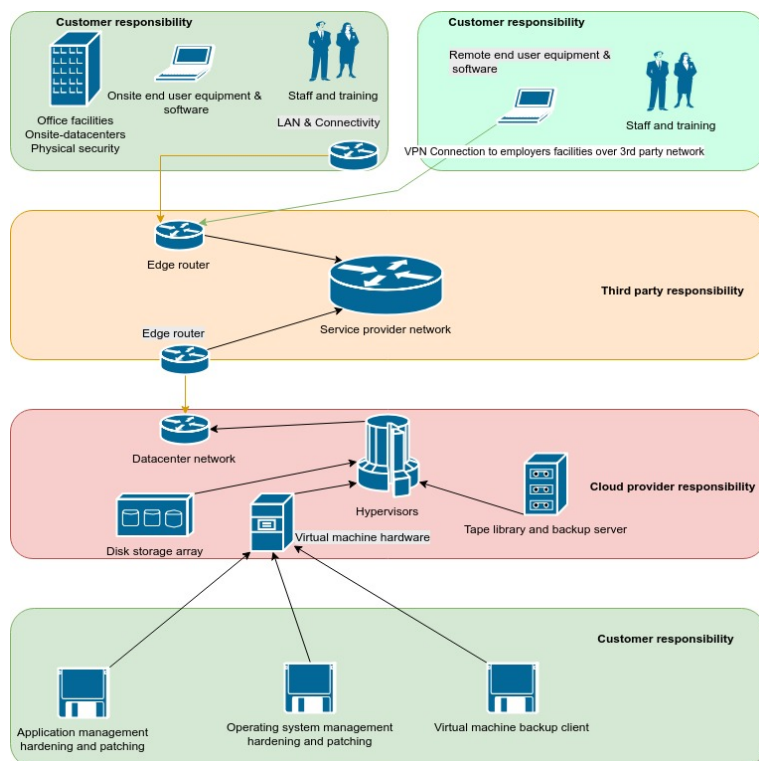


Figure 1: IaaS Architecture and responsibility zones

Given that part of the infrastructure is being outsourced to external service provider it is essential to understand the implications to responsibilities this approach brings with it. This is to illustrate the widening responsibility zones and how the role of connectivity becomes an essential factor for reaching the cloud-hosted services. Also, different failure domains play important role in assessing the feasibility of the IaaS deployment, meaning the failure domain is no longer within the inhouse data center, but it is spread amongst the whole path - both geographical, and otherwise - between the users accessing the service and the servers actually hosting the data. Picture below tries to identify some of the minimum responsibility zones this IaaS creates and how they might map out to a IaaS service. Note that the services illustrated are not part of each and every IaaS service offering, but instead they differ in their content.

2.1.2 Platform-as-a-service, PaaS

As stated above IaaS does not address the many of the scalability issues or automation challenges faced by organizations especially from the perspective of a software. All the parts of the software infrastructure must be provided by the customer. To ease this task PaaS providers can provide software platforms to certain level. Typical software platforms can be for example databases, logging and payments services, which can be used via various APIs (Kavis, Michael. "Platform as a Service". Architecting the Cloud: Design Decisions for Cloud Computing Service Models (SaaS, PaaS, and IaaS). John Wiley & Sons. © 2014. Books24x7. <<http://common.books24x7.com.ezproxy.jamk.fi:2048/toc.aspx?bookid=62597>> (accessed September 25, 2018)) Several PaaS related technologies also aim at automating the provisioning procedures for the virtual machines and containers that actually run the application. Examples of these services could be for example a Kubernetes platform, which would provide an API for containers for automatic scalability. Containers are relatively new concept in computing but they are used to package the application and its dependencies in to a manageable units for distribution and running in cloud platform. (What is a container? Docker documentation 2018). These containers can then be housed in orchestration tools such as the aforementioned Kubernetes or Docker swarm. So, PaaS deployment could be considered being one level above the Software as a Service deployment as it eliminated the need for customer-owned infrastructure for the deployment of a software application. (Saara Suikkanen 2013, s14) Usually PaaS service includes the IaaS services as well, but this is not a given. Many big cloud providers like AWS provide tend to provide a variety of platforms in addition to the plain infrastructure to meet different needs.

2.1.3 Software-as-a-service, SaaS

Software as a service is a method of delivering software application from cloud via Internet connectivity with the least amount of manual work from customer. Using SaaS only requires configuration and user management from the customer, leaving everything else for the service provider. The advantages to the customer is the lack of need to maintain the platform and not needing any personnel to execute the maintenance tasks, which is beneficial especially when talking about services that do not belong to the core functionalities of the customer. (Kavis, Michael. "Software as a Service".

Architecting the Cloud: Design Decisions for Cloud Computing Service Models (SaaS, PaaS, and IaaS). John Wiley & Sons. © 2014. Books24x7. <<http://common.books24x7.com.ezproxy.jamk.fi:2048/toc.aspx?bookid=62597>> (accessed September 25, 2018)) Naturally, the SaaS services can not be used for software that require any heavier tailoring than just predefined configuration changes. A real-life example that illustrates the stacking of cloud services and the SaaS could be a email service that has its customer specific frontends running in containers on service provider orchestration tool that utilizes virtual machines housed in service provider facilities and hypervisors somewhere. SaaS services are nowadays very common (Kavis, Michael. "Software as a Service". Architecting the Cloud: Design Decisions for Cloud Computing Service Models (SaaS, PaaS, and IaaS). John Wiley & Sons. © 2014. Books24x7. <<http://common.books24x7.com.ezproxy.jamk.fi:2048/toc.aspx?bookid=62597>> (accessed September 25, 2018)). Key point to understand with the SaaS model is that customer has no control of the underlying software deployment or the computing infrastructure in SaaS model. (STAMOU, Aikaterini, 2014,s5). This is essential differentiator between SaaS and PaaS.

2.1.4 Differences between cloud types

Private cloud is understood as a service that is being operated by a service provider as a service to be used by single customer. Private cloud tends to use the same techniques as public cloud but they are configured to help the customer organization be more responsive and efficient in the IT resource usage than with traditional IT operation model. (Saara Suikkanen, 2013 s11) There are generally two types of private clouds, ones that are hosted on the customer premises and then there are those that are hosted on service provider infrastructure. It should be noted that while cloud infrastructure could be externally hosted, it is still considered a private cloud if the infrastructure is solely used by single customer organization. (Juha Ahlgren 2012, s10). Infrastructure on public cloud on the otherhand is shared amongst the various customers of a service provider. (TAMOU, Aikaterini, 2014, s6). Public cloud is hosted on the service provider facilities and all maintenance, modifications and upgrades are done by the service provider, meaning that customer has no control of the hosted infrastructure. One exception to this is the fact that certain service providers give customer the options of choosing from several geographical locations from which to run their service. (Juha

Ahlgren 2012, s12). The economy of scale can mean that the public cloud can offer efficient storage, compute and connectivity at reasonable price. This can be especially true with the charging models where customers are required to pay only for the service they require and use. (Saara Suikka, 2013, s11) Combinations of the public and private cloud are called hybrid clouds. These clouds can tie the infrastructures of private and public cloud together and allow the customer to extend their capacity beyond what is available in the private cloud by additionally utilizing the public cloud on time of need. This is called cloud bursting. Meaning that customer uses private cloud under normal circumstances but during peak load some or all parts of the service can be transported to public cloud. (Juha Ahlgren 2012, s13) Fourth and final form of cloud is the community cloud. Community cloud is a multitenant cloud setup that is utilized by several organizations that may share a common interest or computing concerns. Such concern could come in a form of a compliance requirement, audit requirement or that the organizations require high speed access to common data, for example research organizations working on a common project. (Saara Suikkanen 2013, s12) There are two different approaches to the cloud deployment in general, private cloud and public cloud. With private cloud the data and services are hosted hardware dedicated to particular end user organization, while public cloud shares the same hardware infrastructure across all the tenants. This is important aspect to understand when planning application deployment to cloud environment as private cloud provides few key security enhancements per default while at the same time some of the main benefits of cloud diminish:

- CPU and memory is dedicated so computational performance likely more predictable
- Storage performance is easier to predict in terms of IOPS and especially latency
- Vulnerabilities such as recent spectre and meltdown are limited to within one organization
- Network capacity is easier to predict and it is possible to have less latency if housed locally
- In some cases, privately housed and managed cloud might be more likely to pass audit for handling classified information

First three of the above are fairly obvious, but the notion about using privately housed and managed cloud for audit compliance is worth of further clarification. Finnish ministry of finance states

in the VAHTI that when using shared capacity the organization is leaving the management and control of the platform to service provider. This highlights the importance for the client organization to have efficient methods and processes for overseeing the service provider and to understand the nature of the information stored in the cloud and to have efficient means of confirming that the service provider acts according to the requirements and contracts. (VAHTIlohje, vaatimukset tekniselle tietotekniikkaympäristölle, 2016). This could also be understood so that it is quite possible to run classified application in private cloud and gain all the provisioning benefits, as long as the environment supports the required security controls, management processes and that there are otherwise sustainable technical environment to run the application on. This is true regardless if the private cloud is hosted inside the customer company or physically from service provider facility.

Public cloud on the other hand is generally purely shared capacity housed in service provider facilities and usually boasts all the outsourcing benefits within the cloud offering. While some public cloud offerings may have dedicated capacity, it is more often built around the performance factor. Example of this kind of performance oriented offering could be fatnodes with large amounts of RAM or GPU-computation capacity that is rented on hourly basis. That being said, public cloud can provide additional security features that can compensate for the shared infrastructure, like:

- Data encryption at rest as given, with custom encryption keys
- Multifactor authentication to cloud management portal
- Redundant connectivity from service provider facilities

It goes without saying that the list above is not by no means complete and that all of the above are possible to implement in private cloud as well.

2.1.5 Mixtures of the two, aka hybrid cloud

For some implementations it would be ideal to combine the two, private and public cloud. Example of this type of deployment might be an application that has a Internet facing frontend running in service provider cloud, utilizing load balancers and firewalls offered as part of the cloud offering, while the database portion of the application is housed inside the customer datacenter, in customer

hardware. Many vendors boast connectivity services that allow the "stretching" of in-house data-center to service provider facilities. While flexible, this also creates complexity as it is more difficult to understand the dependencies of applications deployed in this manner. Also, it is worth noting that interconnecting data centers is quite complex topic on its own, especially if workloads are to be moved between locations. Some of the risks and implementation models are discussed in the later chapters.

3 Security in cloud

Security in cloud has several characteristics compared to traditional server infrastructure, especially in public and community cloud infrastructure. Private cloud mostly follows the traditional threat scenarios, nevertheless some aspects can be seen even here. In this chapter, we will cover several problems we have seen in cloud environments.

3.1 Security in private cloud

- Trusted environment - Self hosted, or hosted by a trusted partner The biggest difference between traditional server environments or even most of the virtual machine environments is the freedom often provided by the cloud environment. Thus users with less security awareness may be creating and maintaining instances, which may lead to poorer quality of servers in terms of security, a good example is servers meant for testing purposes. Also, if creating servers is more free across the personnel, this might lead to vague responsibilities and neglect of maintenance standards. An organisation has several means to mitigate these problems, for example by providing custom images providing a certain level of security, by running security scans against the servers or by training to name a few examples. (Todo: Generic stuff about backups and their role in security.)For recovery purposes the cloud environments don't necessarily provide decent native backup solutions. This used to be the case for example with OpenStack's earlier version. This easily lead to situations, where backups were easily neglected. Thus if there is a data loss by technical failure or as a result of a system breach, bringing up the system could be quite troublesome.

3.2 Security in community cloud

- Partly trusted - Self hosted or hosted by a trusted or known partner

3.3 Security in public cloud

- Hosted by a known or unknown vendor - You can not normally audit the vendor -> blindness - Server location might change (take Apple's China incident) especially if vendor is an international actor - Contract - Mitigations

3.4 Introduction to security aspects in cloud

3.5 Security aspects of application development in public cloud

- Make application too fast for attackers (cloud native) - The application can be provided from distributed location, same applies to data (which could be broken into meaningless pieces to fight data theft) - Possibilities?

3.6 Introduction to cloud security and application deployment

3.6.1 Source code control in cloud application development

3.6.2 Secure processes for updating application and library code

3.6.3 Secure application administration procedures

- Best practices - Limitation created by cloud environment - Connections over internet

3.7 Security and availability, and what to prepare for

Players in the cloud service provider field tend to be well connected to Internet and to each other (FICIX,2018). This on its own gives service providers some edge when compared to inhouse data center and the connectivity options available. It is also fairly safe to assume that service providers also have more personnel available to run their infrastructure compared to small or medium size company whose business is non-IT-infrastructure.

That being said, especially public cloud hosting with containers should not be seen as a similar virtualization platform as privately maintained in-house datacenter. There are numerous examples of cloud failures that could have been mitigated by understanding the nature of cloud computing and by building the application accordingly, e.g built-in distribution of application resources. Let us next review few of these incidents and discuss some of the remedies that could prevent or lessen the damage caused by these kind of incidents.

3.7.1 File sharing service had authentication bug

On June 2011 well known file sharing service faced an issue where some users could login to accounts without using correct password (dropbox,2011). This led to a situation where the files owned by these accounts were possibly exposed. According to the file sharing service they reacted to this issue by "ending all logged sessions". So not only was there a possibility of information leak, but also innocent users were kicked out of the system to safe guard the data, while the initial step of kicking out users in this sort of situation is not all that dramatic. But what could users do to avoid being denied access to their data, and on the other hand, how to protect their files stored in public cloud?

First option would be to upload only encrypted files, this more or less solves the initial issue of data exposure given that reasonably strong encryption is used. The second threat of being left without access to files due to being kicked out of the service could be avoided by using multiple file sharing services, each with identical data. Now, this might sound laboursome but by using programmatic approach to upload files this replication work could be hidden from the end user.

Similarly, downloading files could be done programmatically hiding the underlying issue with one of the file sharing services. This programmatic approach would also lend itself nicely for the encryption as key handling could be automated to some extent as well.

3.7.2 Outsourced datacenter

To discuss the merits and disadvantages of outsourced data center it is essential to define the meaning of outsourcing in the context of cloud services and data centers. Outsourcing is usually understood as improving the efficiency of business by transferring parts of work that are not seen as the core business to 3rd parties who specialize in those particular tasks, such as power distribution, cabling, IP and ethernet connectivity, rack installations, physical access control etc. (datacenter-journal.com, 2012) Grand idea of outsourcing being that the 3rd party would have better set of skills to do the work in more efficient and reliable manner and possibly by means of volume do it cheaper too. In our context this could mean that company does not see running its own data center and all its dependencies as a core business element, instead company wants to focus on application development and selling application support and maintenance without too many regulations set by authorities or customers. This kind of requirements create decent starting point for investigating if particular application or set of applications could run in outsourced data center, possibly in cloud fashion.

3.7.3 Outsourced server etc. infrastructure, either partially or wholly

As previously stated, there are few levels when it comes to outsourcing. Namely in our context this means either outsourcing all of the infrastructure required to run the application or parts of it. There could be situations where customer has enough resources to run their own servers and networking setup but they might lack the expertise required to maintain storage hardware or HPC connectivity. These kind of mixtures of both outsourced and inhouse platforms require care and good communication with the service provider. As the responsibility is split between customer and the services provider it means that the fluent communication and processes are key as this is especially the case during troubleshooting or when deploying or planning changes to the environment

that will require coordination. This is to avoid blamestorming in case something unexpected happens or when capacity expansions and scaling are required in timely fashion.

3.7.4 Large cloud providers are certified, are they?

3.7.5 Cloud native applications: Maximize the benefits of cloud to protect data and services

Although cloud environments can introduce several risks, they also provide an opportunity to improve the security. Several ideas are introduced in this section. Cloud has the advantage of being versatile, agile and distributed. These features can be used even in the well-known technologies.

Cloud native applications are services, which are specifically build to recognize the surrounding provided by the cloud <todo: find a proper description>. The application can adjust to things like load and location. One idea of fighting attacks with cloud native solutions is provided by Nane Kratzke. With for example open source technologies like Kubernetes or Docker Swarm an environment that would move fast in case an attack is noticed could be made. An example solution could be a container based application, on top of virtual machines that could be in geographically disperse clouds and from different cloud providers. In case there is an attack directed at the platform, the containers could be moved to another environment, for example to another service provider before an attacker can gain a foothold at the environment. According to the paper, there are no known approaches to use this technology for security based solutions (About being the Tortoise or the Hare? A Position Paper on Making Cloud Applications too Fast and Furious for Attackers. Nane Kratzke, 2018. <https://arxiv.org/pdf/1802.03565.pdf>. Read 18.4.2018 <todo: check the date>).

Nevertheless, as availability is considered one aspects of security, this technique is well-known to provide this. To be used in the general context of security, there are some issues with this approach, as it is often hard to recognize underlying attacks. This we believe could be acquired at least to some extend with other traditional tehcnologies like SIEM, log analyzing and bug reports to name a few.

3.7.6 Ready-to-deploy security improvements in cloud

In addition to improving security on application level, it is possible to add to this goal with a hardened cloud platform. The big cloud providers like for example Amazon, Microsoft and Google do provide an extensive lists of measures to improve the security. Mostly these are provided by available security related products. To better understand what products these providers are seeing important, we'll introduce the tools these are providing. Monitoring is provided in different forms. For example Azure has a Azure Security Center, which provides views to different aspects of the cloud. Encryption is a common service to provide for all data, but there are big differences in key management, as in Amazon AWS it is possible to bring your own HMS device to improve the encryptions' security. It is also common to provide DDoS protection, providing this can be naturally also a benefit to the cloud provider itself, as as a multi user environment they are definitely targetted frequently through their customer base. Lastly identity management solutions are provided as an asset to ensure secure authentication and authorization. In addition to these examples, every one of these three providers have several other products, which seem to also derive their target customers from others. (AWS: <https://aws.amazon.com/products/security/>, <https://docs.microsoft.com/en-us/azure/security/>, <https://cloud.google.com/security/products/>)

Sometimes it is possible to affect to the cloud installation, and in this case it becomes possible to create specific hardenings on the infrastructure's configuration. The individual components can be made more secure with additional configuration. It is also possible to add certain services like logging and monitoring, hardened base images, secure network configuration and also by auditing the systems security to name a few examples. From open source cloud platform software providers at least OpenStack provides guides for ensuring the security of the installation (<https://docs.openstack.org/security-guide/>). As security is not just about technology, the administrative procedures and processes can be fine tuned to mitigate many risks.

Even though a customer can add a lot to improve the security, there is only limited data available on how the public cloud providers ensure security in operations. The providers might have an extensive list of compliances like (<https://cloud.google.com/security/compliance/>, AWS artifact), but it is unlikely that even a mid-sized actor could make a proper audit for understandable reasons. Also

the role of the service provider doesn't necessarily fit into the criteria's requirement, as the data owner is still the customer and not the provider. This can make a requirement and its fulfillment even vague. Also things like comprehensive and verbose audit reports from the existing are not available, so there is no data from for example restrictions done during the audit if there is a will to rely on existing audits.

3.8 Security and availability cons of cloud deployments

Security tends to be one of the largest concerns when when it comes to cloud computing and this is easy to understand as outsourced data center with cloud computing almost always exposes clients IT systems to third party such as service provider. (Data center journal, 2012). As discussed in the connectivity dependency chapter, it is essential to have working connectivity to the application from campus, or if the application is solely dependant to Internet it should redundant access from the service provider facilities. While particular the service provider might be certified against ISO 27001 for example, it is still essential to have means of control for the impacts of data center outsourcing, this is especially true if ones application or service has to pass an audit from regulatory officials. To summarize the following cons can be identified:

- Exposure of infrastructure to 3rd party
- Importance of reliable Internet or WAN connectivity
- Requires additional control for 3rd party if subject to regulation

3.8.1 Service level agreements

Service level agreements aka SLAs are sets of condition and terms defined in contracts between customer and the service provider. SLAs can be used to define and agree upon the service levels between provider and customer, including sanctions if the terms are not met. Conditions and terms in SLAs can include various technical, commercial and business service level objectives (SLOs) combined with mechanics of how to measure that the agreed upon services levels are met. (Sta-

mou,2014)

To successfully utilize SLA as a way to improve service availability and security the SLA life-cycle could be split into four parts as follows: (Stamou,2014)

- Creation of the SLA including contract
- Implementing the SLA
- Enforcement and monitoring of the SLA
- Termination of the SLA

Generally speaking the first step consists of service provider predefines set of various SLA levels for customer to choose from and to bound the contract upon. These could be considered as templates for the SLA. Customer then reviews these templates selects one, possibly modifying it and sending it back to service provider for a review. Service provider then accepts, declines or sends modified version to customer for a review. (Stamou,2014 s 13). Rest of the SLA life-cycle consists of implementation, regular reviews and eventually ending of the SLA as stated above.

What makes the SLA for cloud specially tricky is the fact that currently SLAs for cloud lack standardization this is bad as standardization would lead into more structured content of SLAs. In a perfect world the SLA should take into account the individual risk requirements of the customer, this can lead into highly tailored SLAs. (Stamou,2014 s14).

3.8.2 Privacy implications of running application in cloud

3.8.3 Dependant to external 3rd party provider

3.8.4 Dependency to connectivity

Since application is running the remote location it is obvious that connectivity is of paramount importance, in essence, having no connectivity in the campus means having no application and this can mean having no business. While many organizations have Internet connectivity these days it is still surprisingly uncommon for organizations to have backup connectivity if the unthinkable cut

happens. Fiber cuts are not all that uncommon (Teddy Hayford-Acquah and Ben Asante, 2017) and while carrier backbone networks data center connections tend to be redundant in design, the so called "last mile" connection from service provider point-of-presence to customer can be a single point of failure. These non-redundant links can cause catastrophic failures when cut happens during business hours. There are couple of approaches to mitigate the availability issue presented by these failures as discussed next.

To reduce the impact of last mile failure it is common practice to have two physically separate lines from service provider, terminated to two separate customer premisses routers in two separate equipment rooms. (Gunnar Bøe, Vidar Faltinsen, Einar Lillebrygfjeld, 2011) Two routers using VRRP protocol act in active-passive manner to provide so-called first-hop redundancy (rfc5798, 2010). This approach, when combined with physically separate lines, provides protection from fiber cuts on the last mile and this also protects from power supply failures in the customer premisses router and also acts as a backup connection during router software upgrades and some configuration changes. However, this method does not protect against catastrophic failures in the service provider network. To accomplish this, it is required to have similarly separated lines and routers from two separate service providers. Assuming that the customer has some IP block(s) to announce over BGP and that the service providers are accepting the customer IP block(s) for transit it is possible to create fully redundant last mile connectivity. All these relatively complex and expensive requirements are likely the reason why organizations wont purchase redundant connectivity but instead accept the risk of significant business impact and downtime. (Packetworks, 2016).

Similarly, if cloud application is running in a "stretched" network infrastructure, eg. data center interconnect, it is essential that the interconnect is built in redundant fashion. While redundancy is all good it can also cause failures of different kind, but with equally potential for catastrophe (Pepelnjak, 2011) as we'll discuss next.

As most of the applications - in cloud, or otherwise - rely on IP connectivity it is unfortunate misconception that the underlying connectivity from two or more data centers and between datacenters is a given and identical. Having separate IP networks in data centers means that the server, container or whatever running the application loses its connectivity during workload transfer without recon-

figuration of IP address and possibly DNS. Classic solution is to leave out the IP from the equation and simply stretch the OSI layer 2 ethernet via means of OSI layer 1 technology to two or more data-centers. Given the broadcast nature of ethernet and how it behaves during loops this could be seen as a simple but dangerous approach that effectively stretches the data center but at the same time creates a large failure domain, eg. the feat of fate sharing if one data center has a loop the other one goes down as well. This is especially bad for applications that rely on OSI layer 2 signalling, such as redundant routers and firewalls (Pepelnjak, 2011). This problem with location redundancy could be solved by making the application layer not so reliant on the underlying IP layer, this could be done for example by decoupling the service IP - address that end users connect to - and advertising it to data center routers via BGP over only locally significant subnet. Even while there are tools for this sort of decoupling (RIPE 2010), this kind of approaches have apparently been deemed as non-trivial and time consuming tasks so currently it would appear that the accepted solution is to introduce more complexity outside the application to hide the underlying already existing complexity of IP transport. One such method is overlay networking, such as VXLAN, that builds up a stretched OSI layer 2 domain over routed network (rfc7348).

Above has been all about IP applications, yet majority of classical applications require storage to function. This means that storage has to be mirrored or replicated as well. As application generates value of some kind and it needs to store this newly created value somewhere this tends to mean that storage simply has to work. This is one of the reasons that still keeps fibre channel protocol alive, even while ethernet has gained whole load of new features that bring it closer to par with fibre channel when it comes to acting as a transport for storage traffic, eg. lossless connectivity via flow control and in-order-delivery. (Cisco, 2010 and Alex Grossman). Be that as it may, even FCoE still retains the fibre channel provisioning concepts such as zoning, masking and aliases and it is very much OSI layer 2 technology, similarly to native fibre channel. Traditional fibre channel implementations has two, physically and logically separated fabrics to guard against both hardware failures and misconfigurations. (Brocade, 2001). This is could be seen as a overkill these days, but when it comes to safeguarding ones most critical assets of an organization and what has been seen in the wild by the authors, this design still holds truth. Obviously, fibre channel supports stretching of fabrics over geographical distances in highly available manner, within the latency window, but when

it comes to virtualizing multiple tenants into physical fibre channel fabric the amount of scaling is limited (Brocade, 2015). This makes it non trivial for service providers to provide adequately separated, multi tenant storage connectivity from shared infrastructure. Hence, it is likely that if ones application deployment in private cloud relies on service provider facilities to provide on par redundancy for storage replication price tag might be significant. Fibre channel on public cloud is not an option.

Another class of connectivity worth mentioning is the interprocess communication(IPC) and remote direct memory access (RDMA) protocols seen in high performance computing. These applications are very sensitive to latency to begin with and this makes these protocols less likely candidates for stretched fabrics. It should be also noted that protocols such as Infiniband are suited to carry IP and storage traffic in addition to their more commonly seen HPC traffic patterns. Yet, similarly to fibre channel infiniband has limited scalability in terms of multitenancy via number of tenant specific partitions available in the infiniband fabric (Mellanox, 2014). These protocols, just like fibre channel, are non trivial to stretch over large distances in feasible and multitenant manner.

Regardless of the protocol being transported from end user to data center there is likely a need for encrypting the data while in transit. One approach is to implement the required encryption in the application itself by utilising TLS. It has been seen that to avoid complexity in applications there is a push to implement the encryption in the connectivity layer. While there are quite a few methods of implementing encryption in network it should be questioned if it is a sustainable choice to outsource application security to network layer. Implementing encryption using IPSEC (rfc4301) commonly indicates that OSI layer 3 routing should be implemented between data centers, while doing routing is healthy choice for data center interconnect in terms of limiting failure domains it also means that overlay networking is likely required if OSI layer 2 transparency is insisted upon. To implement both OSI layer 2 transparency and encryption one could choose to do encryption on OSI layer 2 via MACSEC (Juniper, 2018), VXLAN over IPSEC or by utilising encryption in DWDM (Arista, 2018) level. It should be noted that both MACSEC and IPSEC have impact in the capacity of performance in terms of payload transferred vs capacity utilized. Running VXLAN over IPSEC may have impact in the net payload as well, or atleast MTU should be carefully considered. Many public cloud providers such

as Amazon (Amazon 2018), Google (Google 2018) and Microsoft (Microsoft 2018) support IPSEC tunnels to tenant specific virtual routing and forwarding instances that are logically separated from one another. Still, it is worth mentioning that even if the data center interconnect from customer data center to cloud provider is encrypted this does not mean that the intra data center traffic inside the service provider facility is encrypted in any fashion. This is one reason why it might be a good idea not to rely on network level to implement the encryption, instead utilize sufficient encryption the application level, just to be sure.

As a conclusion for dependency to connectivity it could be stated that if one's application runs on IP, supports decoupling of service IP from locally relevant IP and it takes care of its own encryption it is likely a good candidate for cloud. Also, campus connectivity from end users to application should be built in fully redundant manner, or at least the risks of non-redundant connectivity should be recognized and accepted at the highest level. Most non-IP services, such as fibre channel and infiniband based applications are difficult to run in a stretched, hybrid cloud and IP based alternatives should be considered.

3.8.5 Application security cannot be outsourced even in SaaS

3.8.6 System security cannot be outsourced in IaaS

4 Self-assessment of cloud security posture

This part of the document describes the questionnaire that acts as a base for self-assessing security of cloud deployment from classical confidentiality, integrity and availability aspects. Questions in the assessment sheet rely on the topical discussions in this paper that can be found from the introduction section. It is to be noted that this is the first version of the assessment and therefore candidate for change.

4.1 What is self-assessment of security and why bother?

The questions below are there to help assessing ones current or upcoming cloud deployment for flaws in the setup and certain caveats that could risk the application deployment. These questions could also be used to challenge the cloud service provider when discussing the proposed deployment model, or when outright selecting services providers via procurement.

4.2 Existing assessment methods and proposed controls: cloud security alliance

There are few assessment sheets and control selections available freely online. Often these documents are suited for specific use cases. Fine example of this is the cloud control matrix by cloud security alliance. (cloudsecurityalliance.org 2018). Aforementioned document is to some extent compatible with at least the following industry-accepted, widely known frameworks: ISO27001/ISO27002, PCI and NIST. This document by cloud security alliance is geared towards checking the compliance of the cloud provider, and on the other hand it is also usable tool to evaluate the maturity level of the controls required by the customer. Document in question is split into sixteen separate sections, each describing topics of their relevant fields, this includes their architectural reference point, corporate governance relevance, applicability to different cloud delivery models and supplier aspects. Next, we'll give a short description of each of these sixteen sections and how we see their relevance for a cloud customer in Finland.

4.2.1 Application and interface security

This section of the cloud security alliance document is about security and compliancy. Document start off by stating that application programming interfaces should be designed, implemented and tested according to some known standard, such as Open web application security project, from here on known as OWASP. It is also stated that before any access to assets is given any contractual and regulatory requirements for customer access should be addressed. Data integrity is discussed from two separate aspects. From technical security viewpoint this involves input and output validation, error detection and correction and mechanics to detect misuse. Policies and procedures

are also discussed from data integrity standpoint as it is stated that policies and procedures shall be documented, established and maintained to support security. This includes the CIA triangle across various system interfaces, jurisdictions, data life cycle management and business functions.

Our take on the aforementioned topics raised by cloud security alliance is that key elements to protect data are mentioned. It is especially important to point out OWASP as an example control framework instead of coming up with ones own set of web application or API security controls. On the other hand, given how many APIs are involved in even the simplest of cloud application deployment it could be noted that the amount of time and effort it takes to make sure that every single API element is in compliance with OWASP is significant factor. To this end, it is good that contractual and regulatory aspects are also mentioned as these could help to enforce it that the required resources are available to fulfil the compliance as they might be seen as hard requirement for the business. While regulatory requirements may not necessarily obligate the service provider side, they do help to guide the customer to select service providers who are more likely to be compliant.

4.2.2 Audit assurance and compliance

This section of the cloud security alliance document talks about auditing and how audit plans need to be developed and maintained in order to support business processes. It is stated that the idea of audit plan is to review the functionality of the implemented security controls. It is also clearly stated that all audit actions should be agreed upon prior to any auditing takes place. Regular reviews and assessments should also take place to make sure that the organization takes up nonformities or established policies, standards, procedures and compliance obligations. These things should be maintained in a control framework that captures key requirements, both contractual and regulatory, and this control framework is to be reviewed regularly.

To us the key take away from the aforementioned is the regular review and updating of the control frameworks. This can be illustrated by the recent fuzz caused by the general data protection regulation, from here on GDPR. GDPR is good example of situation where pre-existing frameworks might have required even considerable renewing to match the new requirements forced by regulator. An-

other essential point made in the text is the usage of outside auditor as this can greatly help the organization in their efforts of avoiding becoming stagnant and also provide external view to the internal processes and controls. Auditing can also be seen as part of product development cycle, as something that helps to improve quality and quality control.

4.2.3 Business continuity management & operational resilience

This section of the document largely describes, as expected, various areas of business continuity and their domains. Examples of these could be implementing the very idea of business continuity planning and making the organization aware of what and why this documentation is required. This includes building of the basic requirements for business continuity planning namely defining the scope and purpose of the plan, who should use the plan, who owns it, describing the communication parties and channels required to implement the plan and last, but by no means least, method to invoice, cancel and eventually stop the active business continuity plan activities. It is also stated that the aforementioned business continuity plan shall be tested, both in theory and in practice as the need may be. This testing could include scenarios such as environmental hazards, technical cyber security risks and organizational changes. This testing should also take place on regular interval. Document also states the importance of up-to-date documentation of systems and services that run on those systems. This includes things like administrator documentation, end user documentation and architecture diagrams. These documents should be available to the parties involved in operations of the services and systems. It is said in the document that there should be a standardized way to measure the impact of the hazards. To this end, it is essential to identify the critical components that make up the service, including their dependencies.

In our eyes these are essential requirements. Lack of the aforementioned documentation is a warning sign of maturity level of the service provider. While the requirements stated here are no doubt important, they also present a significant requirement for resources both human and monetary. Hence, we'd like to emphasize the importance of commitment from top level management that the required resources are made available, as otherwise these documents will be just that, documents. While regular review of the policies and plans was mentioned, we feel that this part should

be emphasized even more as out-dated document can be dangerous and on the otherhand situations with things like environment and hardware vendors might change and the plans should reflect these changes. These reviews should be part of yearclock. These documents might not only be seen as internal documents, but also also as basis of producing extra value by making it easier to interface with customers. This could mean things like confirming the customers also has similar documentation, and possibly aiding the customer in creation of such documents based on services providers pre-existing documentation. This shows commitment to security, availability and integrity of customer data and it can also be seen as a sign of good will.

4.2.4 Change control & configuration management

These criteria essentially state that there should be a standardized way to interact with vendors, when acquiring new data, applications and acquisitions. Another point is made about how external 3rd parties should also comply with the requirements of the customer, be it testing, release engineering or supply chain management. This section also mentions the requirement of following defined quality control procedures, including established baselines to compare against, this means that only authorized devices, components and software shall be used to implement the change.

We feel that catalogue of authorized components and devices should be emphasized more, while this catalogue should still be kept loose enough to allow innovation. Also, we'd like to point out that the policies for change management should be reviewed and tested regularly, for example how the change management process reacts to unauthorized components possibly found inside the product? This chapter is largely tied with the business continuity management & operational resilience above.

4.3 Risk management in cloud landscape

Risk management is a crucial point in any information systems. Usually a risk is defined as a likelihood of some incident and its consequence to an asset. The risks value is then calculated from

these values and counter measures are decided based on these values. Our opinion is that these standard risk calculations tend to forget one important factor: how much power the actors that are potentially interested in the asset do have. This power can be for example money, time, legislation or law related impact, skill related possibilities like the skills of the actor or the actor's ability to acquire such skills. Forgetting this factor might easily lead into misscategorization of the risk's value. Say there is a company which would like to use commonly used cloud services like Microsoft's O365 and use this as company's collaboration tool with a clause that no classified data can be added to the service. Quite naturally this would categorize the value of the asset quite low. The company doesn't have knowledge of any breaches targeted at the service, thus the likelihood also is very low. But what if the asset interests for example a multi billion actor's interests? Say now this service provides a centralized publicly available target to acquire data like personal data, which again could be used as a source for other types of attacks. Thus the likelihood is kind of higher as it interests an actor with motivation and assets to acquire it. (Stølen, Ketil, Atle Refsdal, and Bjørnar Solhaug. *Cyber-Risk Management*. Springer. © 2015. Books24x7. <<http://common.books24x7.com.ezproxy.jamk.fi:2048/toc.aspx?bookid=115793>> (accessed September 22, 2018))

A cloud based environment presents several new risks, not forgetting the fact it also mitigates some others. Especially with public clouds the customer has only limited capabilities to affect business related decisions, which might pose a great risk. There are examples where actions like this have realized. For example, in February 24 2018 the news reported that Apple had made a decision to move the encryption keys of Chinese iCloud customers to a Chinese service provider. This was presumed to increase the risk of a governmental actor to get access to the keys (<https://www.reuters.com/article/us-china-apple-icloud-insight/apple-moves-to-store-icloud-keys-in-china-raising-human-rights-fears-idUSKCN1G8060>, accessed 22.9.2018). This article is a good reminder that there might be providers who would grant access to data under a pressure or legislative force for powerful actors like governments.

Also other very visible risks appear. As long as you have a clearly defined physical boundaries where the data is located, things are a bit simpler. Introducing a geographically separated physical area expands or even breaks this physical boundary. You should now be able to also monitor this separate

location, as well as the interfaces it has to the internal zones.

In addition, risk analysis in a cloud based environment can be harder, when speaking of public clouds. Things like interviews can be a valuable tool when identifying risks (Stølen, Ketil, Atle Refsdal, and Bjørnar Solhaug. *Cyber-Risk Management*. Springer. © 2015. Books24x7. <<http://common.books24x7.com.ezproxy.jamk.fi:2048/toc.aspx?bookid=115793>> (accessed September 22, 2018)), but interviewing any technical persons from a third party hosted cloud can be a real challenge. This can lead into a crippled risk identification process if not taken properly into account. It is obvious that this type of problem can realise even in non-cloud environments where the personnel isn't committed to the process.

4.4 Self-assessment questionnaire

4.4.1 Main points of interest and reasoning behind assessment questions

5 Auditing application deployment in cloud

5.1 Classified data: to cloud or not?

5.2 Overview of audit criteria for cloud environment

6 Conclusions

6.1 Complexity and applicability of audit criteria

6.2 Review of assessment and criteria

6.3 Further developments

7 References

Lilia Sampaio, Fábio Silva, Amanda Souza, Andrey Brito, Pascal Felber Secure and Privacy-Aware Data Dissemination March 2018

Pablo Chico de Guzman, Felipe Gorostiaga, Cesar Sanchez i2kit: A Tool for Immutable Infrastructure Deployments based on Lightweight Virtual Machines specialized to run Containers Feb 2018

Nane Kratzke About being the Tortoise or the Hare? A Position Paper on Making Cloud Applications too Fast and Furious for Attackers Feb 2018

V. Raisanen, G. Grotefeld, A. Morton Network performance measurement with periodic streams Nov 2002

Finnish Communication Regulatory Authority Pilvipalveluiden turvallisuus - Mitä organisaatioiden tulisi huomioida pilvipalveluja hyödyntäessä

Finnish Communication Regulatory Authority Ohje 5/2014 Pilvipalveluiden turvallisuus May 2014

Microsoft Azure What is cloud computing? - A beginner's guide Published on <https://azure.microsoft.com/en-in/overview/what-is-cloud-computing/>

Docker.com What is a container? Published on <https://www.docker.com/what-container>

Finnish ministry of finance Vahtiohje - Vaatimukset tekniselle tietotekniikkaympäristölle. Published on <https://www.vahtiohje.fi/web/guest/vaatimukset-tekniselle-tietotekniikkaymparistolle>

Causes of Fiber Cut and the Recommendation to Solve the Problem Teddy Hayford-Acquah Ben Asante 2017

Distributed firewalls: how badly do you want to fail? Ivan Pepelnjak 2011

GEANT Recommendations for a redundant campus network Gunnar Bøe Vidar Faltinsen Einar Lillebryggfeld 2011

Packetworks Published on <http://www.packetworks.net/blog/the-risks-of-not-having-business-internet-redundancy.htm> 2016

IETF VXLAN M. Mahalingam D. Dutt K. Duda P. Agarwal L. Kreeger T. Sridhar M. Bursell C. Wright August 2014

IETF VRRP S. Nadas

Cisco Fibre Channel over Ethernet 2010

Brocade Brocade SAN Scalability Guidelines 2015

Alex Grossman Fibre Channel vs Ethernet Published online at <http://www.keycodemedia.com/2017/01/01/fibre-channel-vs-ethernet-by-alex-grossman/> on 4th of April, 2018

Mellanox Switch-IB EDR Switch Silicon 2014

RIPE NCC Thomas Mangin ExaBGP - A new Tool to Interact with BGP 2010

Juniper 2018

Arista 2016

IETF rfc4301 - Security Architecture for the Internet Protocol S. Kent K. Seo

Metropolia Juha Ahlgren 2012

Saimaa University of Applied Sciences Saara Suikkanen 2013

University of Geneva Stamou, Aikaterini 2014

Amazon AWS Managed VPN Connections Available in https://docs.aws.amazon.com/Amazon-VPC/latest/UserGuide/VPC_VPN.html on fourth of April 2018

Google Cloud VPN Overview Available in <https://cloud.google.com/vpn/docs/concepts/overview> on fourth of April 2018

Microsoft Connect your datacenter to Azure Available in <https://azure.microsoft.com/en-us/services/vpn-gateway/> on fourth of April 2018

Datacenter Journal Jeff Clark 2012

FICIX ry FICIX statistics Available from <https://stats-ficix.basen.com/#/page?name=StatsWelcome&source=wiki> on fourth of April 2018